

EXPERIMENT 1

CODE:

```
Import libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score


Load California Housing dataset

california = fetch_california_housing(as_frame=True)
data = california.frame


Display first few rows and dataset information

print("First five rows of the dataset:")
print(data.head())


print("\nDataset Info:")
print(data.info())


Check for missing values

print("\nMissing values in the dataset:")
print(data.isnull().sum())
```

OUTPUT:

First five rows of the dataset:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	\
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	

	Longitude	MedHouseVal
0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 20640 entries, 0 to 20639

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	MedInc	20640 non-null	float64
1	HouseAge	20640 non-null	float64
2	AveRooms	20640 non-null	float64
3	AveBedrms	20640 non-null	float64
4	Population	20640 non-null	float64
5	AveOccup	20640 non-null	float64
6	Latitude	20640 non-null	float64
7	Longitude	20640 non-null	float64
8	MedHouseVal	20640 non-null	float64

dtypes: float64(9)

memory usage: 1.4 MB

None

Missing values in the dataset:

MedInc	0
HouseAge	0
AveRooms	0
AveBedrms	0
Population	0
AveOccup	0
Latitude	0
Longitude	0
MedHouseVal	0
dtype:	int64

EXPLANATION:

1. Loading the California Housing Dataset

- The dataset is fetched using the `fetch_california_housing` function from scikit-learn's dataset module. This dataset is widely used for regression tasks and includes features like median income, average number of rooms, and median house value (target variable).

- The `as_frame=True` argument ensures the dataset is loaded as a pandas DataFrame, making it easy to work with using familiar pandas functions.

2. Exploring the Dataset

- The `data.head()` function displays the first five rows of the dataset. This provides an initial view of the dataset's structure, column names, and the type of data available.

- The dataset includes columns like `MedInc` (median income), `HouseAge` (age of the house), and `MedHouseVal` (median house value). These features will be used as inputs to model the target variable, `MedHouseVal`.

- The `data.info()` function gives insights into the dataset's columns, data types, and non-null counts. It helps confirm that all columns are of numeric type and that there are no missing values.

3. Checking for Missing Values

- The `data.isnull().sum()` function is used to check if there are any missing values in the dataset. In this case, all columns have zero missing values, so no preprocessing steps like imputation or removal are needed.

Summary

- The dataset appears clean and ready for regression modeling. The next steps typically involve preprocessing (like scaling or encoding categorical variables), splitting the data into training and testing sets, and building a regression model to predict `MedHouseVal`.

- The information from this initial exploration helps identify potential features (`MedInc`, `HouseAge`, etc.) and the target variable (`MedHouseVal`) that could be used to build a predictive model.

CODE:

Basic statistics

```
print("\nDataset Summary Statistics:")
```

```
print(data.describe())
```

Heatmap to visualize correlations

```
plt.figure(figsize=(10, 8))  
  
sns.heatmap(data.corr(), annot=True, cmap='coolwarm')  
  
plt.title("Correlation Heatmap")  
  
plt.show()
```

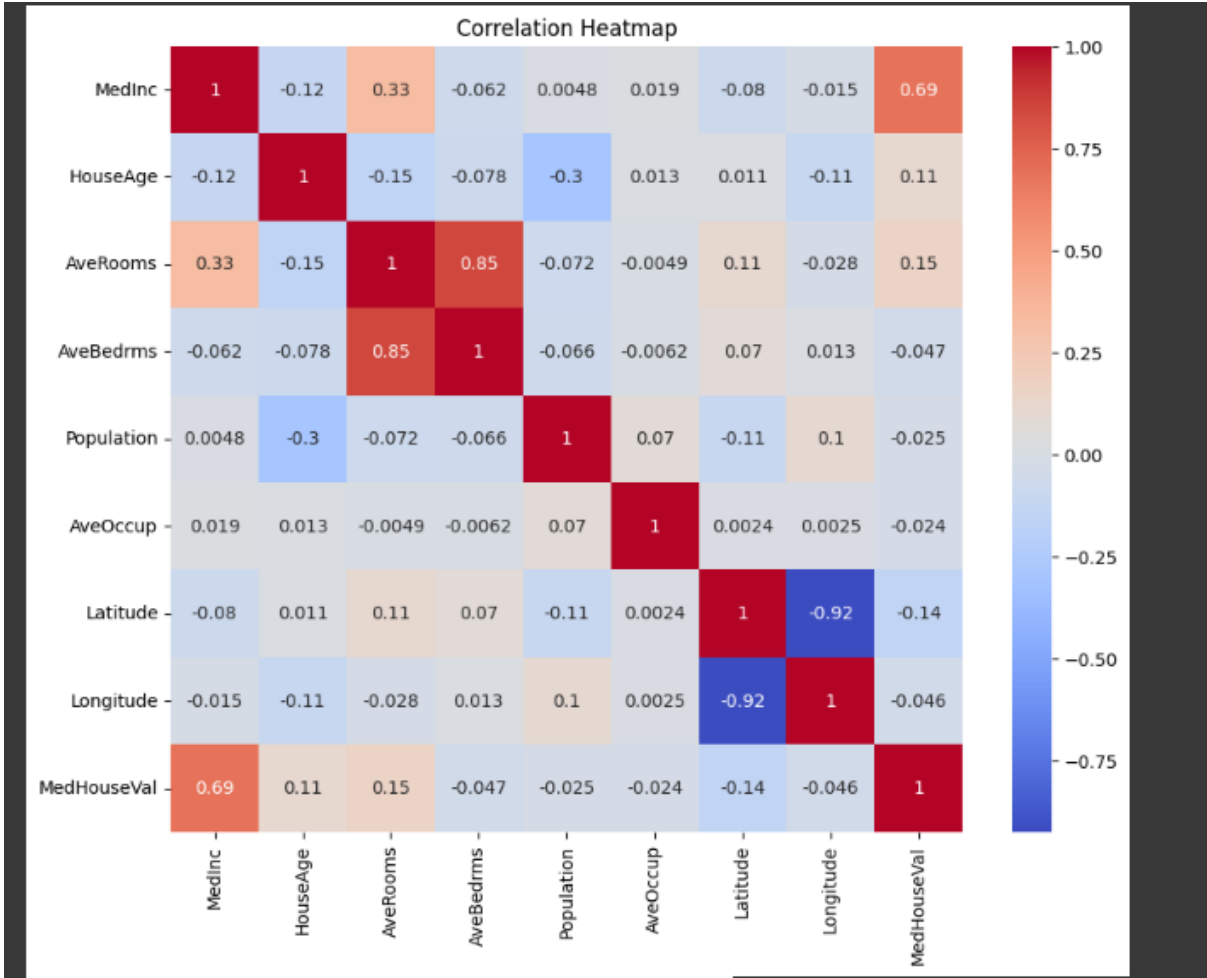
Scatter plot: Median Income vs Median House Value

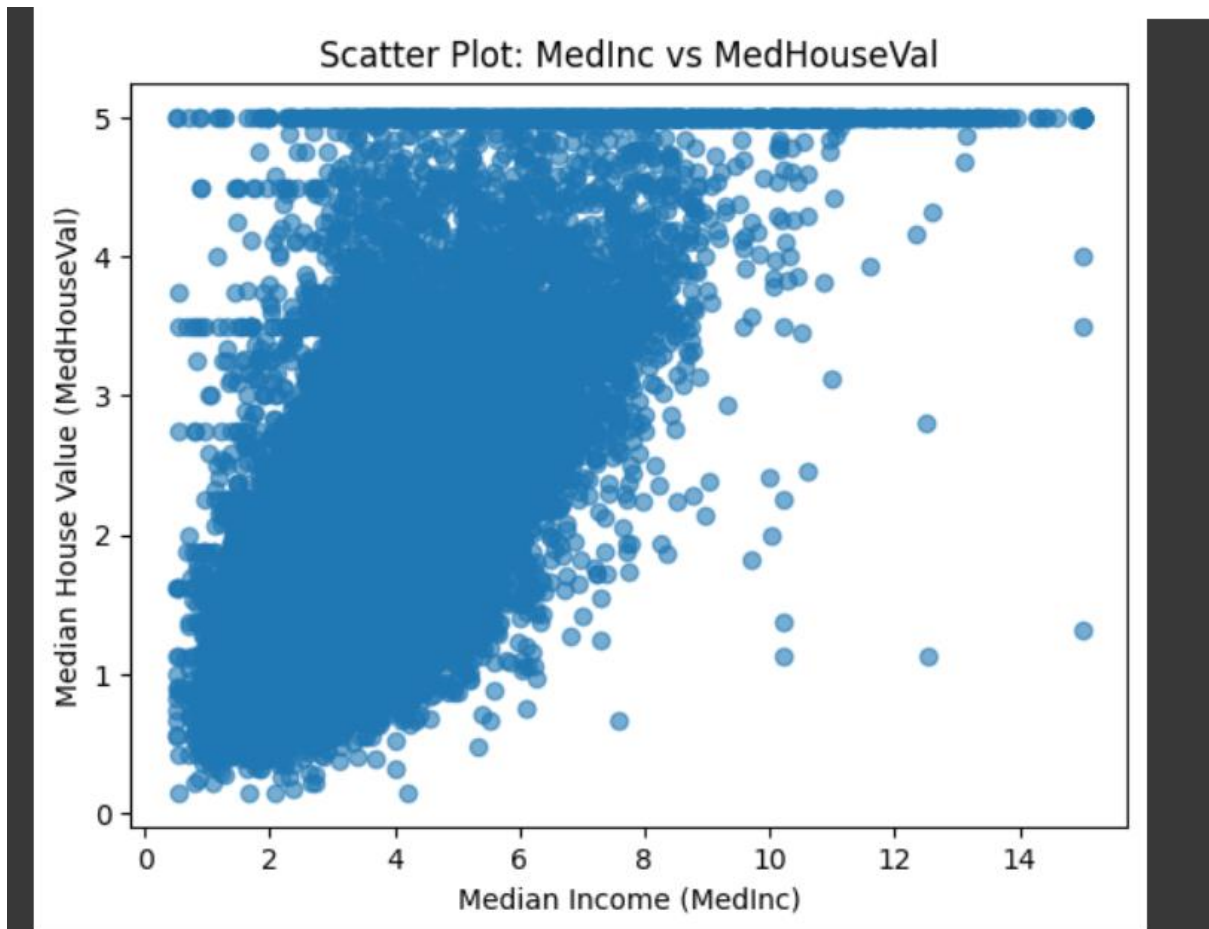
```
plt.scatter(data['MedInc'], data['MedHouseVal'], alpha=0.6)  
  
plt.title("Scatter Plot: MedInc vs MedHouseVal")  
  
plt.xlabel("Median Income (MedInc)")  
  
plt.ylabel("Median House Value (MedHouseVal)")  
  
plt.show()
```

OUTPUT:

Dataset Summary Statistics:					
	MedInc	HouseAge	AveRooms	AveBedrms	Population \
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.870671	28.639486	5.429000	1.096675	1425.476744
std	1.899822	12.585558	2.474173	0.473911	1132.462122
min	0.499900	1.000000	0.846154	0.333333	3.000000
25%	2.563400	18.000000	4.440716	1.006079	787.000000
50%	3.534800	29.000000	5.229129	1.048780	1166.000000
75%	4.743250	37.000000	6.052381	1.099526	1725.000000
max	15.000100	52.000000	141.909091	34.066667	35682.000000

	AveOccup	Latitude	Longitude	MedHouseVal
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.070655	35.631861	-119.569704	2.068558
std	10.386050	2.135952	2.003532	1.153956
min	0.692308	32.540000	-124.350000	0.149990
25%	2.429741	33.930000	-121.800000	1.196000
50%	2.818116	34.260000	-118.490000	1.797000
75%	3.282261	37.710000	-118.010000	2.647250
max	1243.333333	41.950000	-114.310000	5.000010





EXPLANATION:

1. Basic Statistics Summary

- The `data.describe()` function generates basic statistical metrics for each column in the dataset, including:
 - Count: The number of non-null values in each column.
 - Mean: The average value for each feature.
 - Standard Deviation (std): Measures the spread of the values in each column.
 - Min and Max: The minimum and maximum values in the dataset for each feature.
 - Percentiles (25th, 50th, 75th): These provide insights into the distribution of data, showing where the middle 50% of values lie.

This summary is helpful to understand the general distribution and range of the data, identify potential outliers, and decide if any feature scaling or normalization is necessary.

2. Correlation Heatmap

- A heatmap is created to visualize the correlation between different features in the dataset. Correlation indicates the strength and direction of the relationship between pairs of features, ranging from -1 (perfect negative correlation) to +1 (perfect positive correlation). A correlation of 0 means no linear relationship.

- Positive correlations (e.g., between `MedInc` and `MedHouseVal`) indicate that as one feature increases, the other tends to increase as well.

- Negative correlations suggest that as one feature increases, the other tends to decrease.

- No correlation implies that changes in one feature do not systematically relate to changes in another.

The heatmap uses color coding (e.g., cool colors for negative correlations and warm colors for positive correlations) to make these relationships easier to identify at a glance.

3. Scatter Plot: Median Income vs Median House Value

- A scatter plot is used to visually explore the relationship between Median Income (`MedInc`) and Median House Value (`MedHouseVal`).

- Each point on the plot represents a data entry, with the x-axis showing the median income for a region and the y-axis showing the median house value.

- The plot helps identify any trends or patterns between these two variables, such as:

- A positive correlation suggests that regions with higher median incomes tend to have higher median house values.

- Any potential outliers can also be identified on the plot as points that deviate significantly from the general trend.

By using the scatter plot, you get a clearer visual understanding of how these two features relate to each other, and whether a linear regression model might be appropriate to model their relationship.

Summary

- Basic statistics provide an overview of the data's central tendencies, variability, and distribution.

- The correlation heatmap reveals the strength of relationships between various features, which can inform feature selection for modeling.

- The scatter plot offers a visual representation of the relationship between `MedInc` and `MedHouseVal`, helping to confirm any trends or patterns that might inform the choice of modeling techniques.

CODE :

Feature selection

`X = data[['MedInc']]` Using Median Income as the independent variable

`y = data['MedHouseVal']` Median House Value as the target variable

Split the data into training and testing sets

`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)`

Train a Linear Regression model

`model = LinearRegression()`

`model.fit(X_train, y_train)`

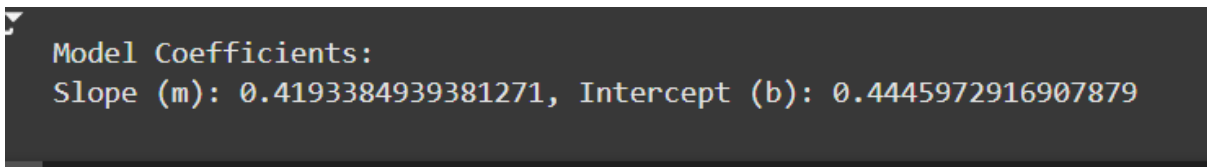
Make predictions

`y_pred = model.predict(X_test)`

`print("\nModel Coefficients:")`

`print(f"Slope (m): {model.coef_[0]}, Intercept (b): {model.intercept_}")`

OUTPUT:

A terminal window with a dark background and light-colored text. It displays the output of the Python code: 'Model Coefficients:' followed by 'Slope (m): 0.4193384939381271, Intercept (b): 0.4445972916907879'.

```
Model Coefficients:  
Slope (m): 0.4193384939381271, Intercept (b): 0.4445972916907879
```

EXPLANATION:

1. Feature Selection

- The feature selection step involves choosing which variables (features) will be used as inputs (independent variables) for the regression model. In this case, only Median Income (`MedInc`) is selected as the feature (independent variable), and Median House Value (`MedHouseVal`) is selected as the target (dependent variable). This simplifies the model to a univariate regression, focusing on the relationship between just these two variables.

2. Splitting the Data

- The dataset is divided into two parts:

- Training Set: 80% of the data is used to train the model, which means that the model will learn the relationship between `MedInc` and `MedHouseVal` based on this subset.

- Testing Set: 20% of the data is reserved for testing the model's performance. This ensures that the model is evaluated on unseen data, which provides a more accurate assessment of its predictive ability.

- The `train_test_split` function performs this division, and the `random_state=42` ensures that the split is reproducible (i.e., the same split will occur each time the code is run).

3. Training the Linear Regression Model

- A Linear Regression model is created using the `LinearRegression()` class from scikit-learn. This model will attempt to find the best-fit line (or hyperplane) that minimizes the difference between predicted and actual values of `MedHouseVal` for the training data.

- The model learns the parameters (slope and intercept) that define the linear relationship between `MedInc` and `MedHouseVal`.

4. Making Predictions

- After training, the model is used to make predictions on the testing set (`X_test`). The predicted values (`y_pred`) are the model's estimates for `MedHouseVal` based on the input values of `MedInc` in the test set.

5. Model Coefficients

- The coefficients of the linear regression model are printed out:

- Slope (m): This represents the change in the target variable (`MedHouseVal`) for each unit increase in the independent variable (`MedInc`). A positive slope indicates that as median income increases, median house value also tends to increase.

- Intercept (b): This is the value of `MedHouseVal` when `MedInc` is zero. In practical terms, it may not always be meaningful (e.g., `MedInc` being zero might not make sense in the context of housing data), but it's part of the equation of the best-fit line.

Summary

- This step involves selecting the independent feature (`MedInc`) and the dependent target (`MedHouseVal`), splitting the data into training and testing sets, and then training a linear regression model to predict house values based on income levels.

- The model coefficients (slope and intercept) provide a mathematical understanding of how median income is expected to influence median house value in the dataset.

CODE :

Calculate metrics

```
mae = mean_absolute_error(y_test, y_pred)
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
r2 = r2_score(y_test, y_pred)
```

```
print("\nEvaluation Metrics:")
```

```
print(f"Mean Absolute Error (MAE): {mae:.2f}")
```

```
print(f"Mean Squared Error (MSE): {mse:.2f}")
```

```
print(f"R-squared (R2): {r2:.2f}")
```

OUTPUT:

```
Evaluation Metrics:  
Mean Absolute Error (MAE): 0.63  
Mean Squared Error (MSE): 0.71  
R-squared (R2): 0.46
```

EXPLANATION:

1. Mean Absolute Error (MAE)

- MAE measures the average absolute difference between the actual values (`y_test`) and the predicted values (`y_pred`). It provides a straightforward interpretation of the model's error:

- Formula: $MAE = (1/n) \sum |y_i - \hat{y}_i|$

- Where `yi` is the actual value, and `ŷi` is the predicted value for each data point.

- A lower MAE indicates better model performance. It's easy to interpret since it represents the average amount by which the predictions are off, in the same units as the target variable (in this case, house values).

2. Mean Squared Error (MSE)

- MSE calculates the average squared difference between the actual and predicted values. Squaring the differences gives more weight to larger errors, making MSE more sensitive to outliers.

- Formula: $MSE = (1/n) \sum (y_i - \hat{y}_i)^2$

- Where `yi` is the actual value, and `ŷi` is the predicted value for each data point.

- MSE is widely used but can be more difficult to interpret directly because the error is in squared units. It penalizes large errors more significantly than MAE, so if there are outliers, the MSE will be higher.

3. R-squared (R^2)

- R-squared is a measure of how well the model explains the variance in the target variable. It tells us the proportion of the variance in the target variable that is predictable from the independent variables.

- Formula: $R^2 = 1 - (\sum(y_i - \hat{y}_i)^2 / \sum(y_i - \bar{y})^2)$

- Where y_i is the actual value, \hat{y}_i is the predicted value, and \bar{y} is the mean of the actual values.

- R^2 values range from 0 to 1:

- 0 means the model does not explain any of the variance in the target variable.

- 1 means the model perfectly explains the variance in the target variable.

- An R^2 closer to 1 indicates that the model does a good job of explaining the relationship between the independent and dependent variables.

Summary

- MAE gives a clear indication of average prediction error in the original units of the target variable.

- MSE penalizes larger errors more than MAE, making it more sensitive to outliers.

- R^2 provides a measure of the model's goodness of fit, indicating how well the model explains the variability in the target variable.

- By examining these metrics, you can evaluate the accuracy and reliability of your linear regression model. A good model will have low MAE and MSE values and an R^2 value close to 1.

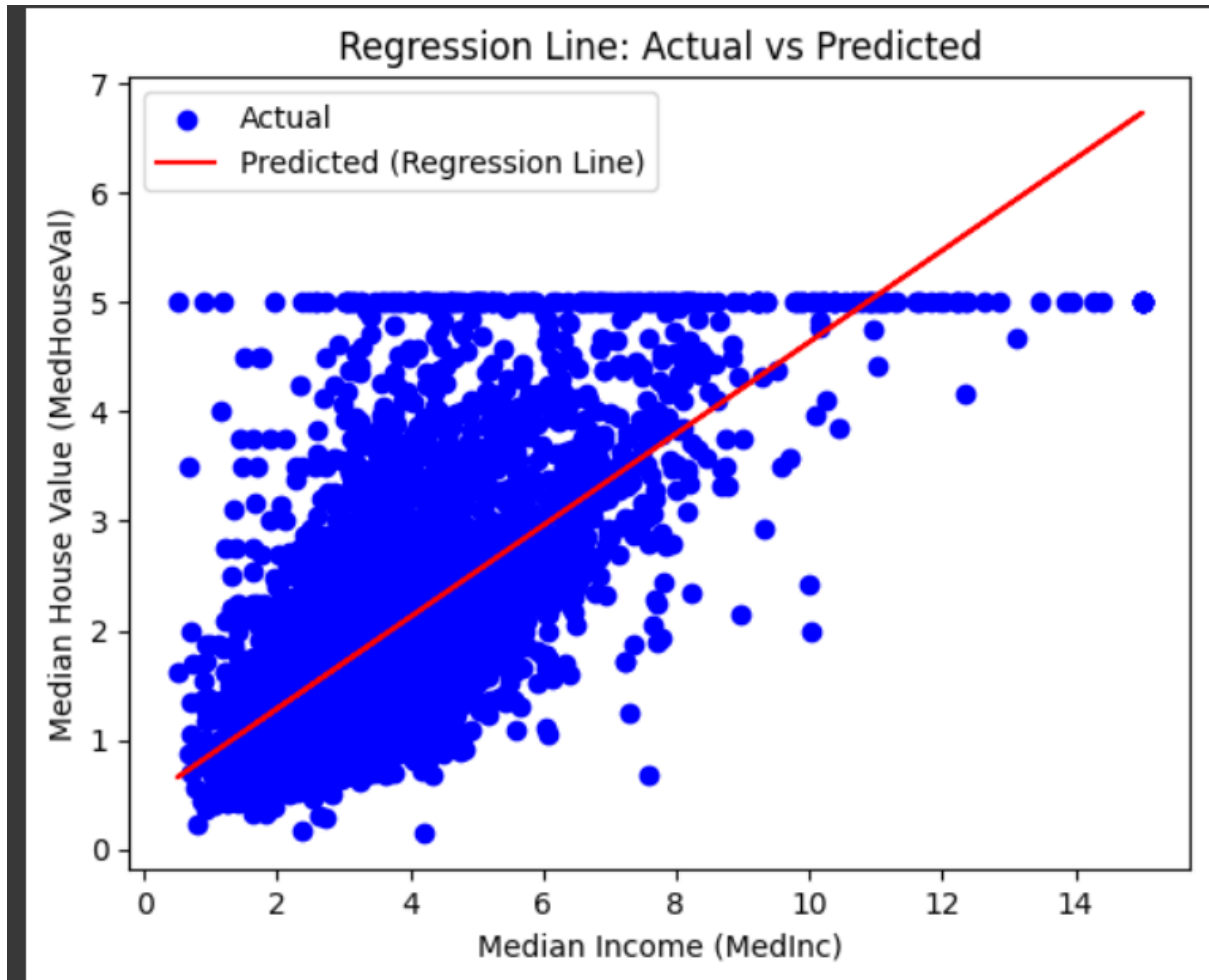
CODE :

Regression line visualization

```
plt.scatter(X_test, y_test, color='blue', label='Actual')
plt.plot(X_test, y_pred, color='red', label='Predicted (Regression Line)')
plt.title("Regression Line: Actual vs Predicted")
plt.xlabel("Median Income (MedInc)")
plt.ylabel("Median House Value (MedHouseVal)")
plt.legend()
```

```
plt.show()
```

OUTPUT:



EXPLANATION:

Visualization of Actual vs Predicted Values

- The plot shows the relationship between the actual values of the target variable (`MedHouseVal`) and the predicted values (`y_pred`) produced by the linear regression model. It provides a visual representation of how well the model fits the data.

1. Scatter Plot of Actual Values

- The scatter plot displays the actual values of the target variable (`MedHouseVal`) from the test set (`X_test` and `y_test`). Each point on the scatter plot represents a data point, with the x-axis showing the median income (`MedInc`) and the y-axis showing the actual house value (`MedHouseVal`).

- This part of the plot helps us see the distribution of actual house values in relation to median income.

2. Plot of Predicted Regression Line

- The red line represents the predicted values of the target variable ('MedHouseVal') based on the model's predictions ('y_pred') for the test set. This line is the best-fit regression line found by the linear regression model.

- The line illustrates the relationship between median income and house value as estimated by the model. It is essentially the linear equation (slope and intercept) derived from the training data.

3. Comparison Between Actual and Predicted

- The scatter points represent the actual data points, while the red regression line represents the model's predictions. If the model is a perfect fit, the red line would pass exactly through all the blue points (actual values).

- The closer the red line is to the blue points, the better the model's predictions are. If the points are spread out significantly from the line, it indicates the model has more error in predicting those values.

Summary

- The plot provides a visual check of how well the linear regression model is capturing the relationship between median income and median house value. The closer the predicted regression line is to the actual data points, the better the model's performance.

- This type of visualization is crucial for assessing whether the linear model is a good fit for the data, helping to identify any significant discrepancies between predicted and actual values.